

User guide

Everything you need to assess an AI use case with this kit. The first one takes about ten minutes. There is nothing to install, download or upload.

1. Assess a use case in three steps

1. **Add a use case.** On the home page, choose *Start your first assessment*. Give the use case a name and an accountable owner; the other details are optional.
2. **Answer the quick check.** It opens by asking whether the solution uses AI when it runs; if it doesn't (for example, it was only built with AI coding tools), AI governance controls don't apply and the kit gives general software pointers instead. Then come eight multiple-choice questions about impact, data, who is affected, the type of AI, autonomy, oversight, access and reversibility. Each answer shows the points it adds, and these eight questions alone set the tier.
3. **Read the result.** You get a risk tier with the reasons behind it, then the controls that apply under the frameworks you selected, then the exports.

GCC AI Governance Kit Guide How scoring works Templates About the data GitHub عربي

OPEN · BILINGUAL · BROWSER-ONLY

Adopt AI governance in one afternoon.

Register an AI use case, get an explainable risk tier, and export the controls that apply under the frameworks you follow. Everything stays in your browser.

1 Add a use case
Give it a name and an accountable owner.

2 Answer 8 questions
About two minutes, right here in your browser. There is nothing to download or upload.

3 Get your tier, controls and PDF
An explainable risk tier, the controls that apply, and a report for your committee.

+ Start your first assessment

Frameworks you follow

☒ UAE AI Charter ☒ UAE Gov Services Code ☒ Dubai AI Ethics ☒ SDAIA AI Ethics ☒ ISO/IEC 42001 ☒ NIST AI RMF ☒ EU AI Act

Show All Save or restore

Your register is empty. Add your first AI use case, or load two examples to see how it works.

The home page: three steps and one button to start.

Choose the frameworks you follow on the home page. The control list and the exports only show references to those frameworks.

GCC AI Governance Kit

[Guide](#) [How scoring works](#) [Templates](#) [About the data](#) [GitHub](#) [عربي](#)

[Register](#) / ▲ High risk Citizen service triage agent

✓ Details

2 Quick check

✓ Risk tier

4 Deep-dive

5 Controls

6 Export

Quick check

Does the solution use AI when it runs?

Choose No if AI was only used to build it (for example an AI coding assistant) and the running product has no AI model or AI service inside.

► Commercial or B2B tool? How to answer

☒ Yes: it uses an AI model or AI service when it runs

☐ No: AI was only used to build it, or not at all

Eight questions, about two minutes. **8/8**

1. What does the AI's output influence?

Think about the most serious thing the output can lead to, not the usual case.

► Commercial or B2B tool? How to answer

☐ Internal productivity only (drafts, summaries, search) +0

☐ Information given to the public or customers +5

The quick check: eight questions, about two minutes.

For a fuller picture, complete the deep-dive: ten more questions plus three optional prioritisation questions. The deep-dive can raise the tier but never lower it. Complete it before any High-risk use case goes to committee.

GCC AI Governance Kit

[Guide](#) [How scoring works](#) [Templates](#) [About the data](#) [GitHub](#) [عربي](#)

[Register](#) / ▲ High risk Citizen service triage agent

✓ Details

✓ Quick check

3 Risk tier

4 Deep-dive

5 Controls

6 Export

Risk tier

▲ High risk

Your tier comes from the 8 quick-check questions. The optional sections only refine it or recommend autonomy: skipping them never makes the tier higher.

Why this tier

- Recommendations to staff about specific people
- It acts on its own; people monitor
- Writes to systems of record, makes payments, or sends external messages

▼ Points breakdown (78 points)

Question	points
What does the AI's output influence?	15
What data does it use?	10
Who could be affected?	8

The result: a tier, the top reasons and the full points breakdown.

2. What the risk tiers mean

The tier is the higher of two results: the points from your answers, and any named rule that sets a minimum tier. [How scoring works](#) lists every weight.

What would lower your tier. Once the deep-dive is done, the tier card shows the tier you would reach by closing operational gaps (testing, monitoring, incident handling, hosting and disclosure), with the points each fix saves. If the tier would not move, the cause is structural (what the AI does, its data, autonomy and access), and only a design change would lower it.

● Little or no risk

Under 20 points. Keep it on the register and apply the baseline controls the kit lists.

◆ Limited risk

20 to 44 points. More controls apply; review them with the accountable owner before launch.

▲ High risk

45 points or more, or a named rule applies (for example, the AI makes decisions about people without a person approving each one). The full control set applies. Complete the deep-dive and take the committee report to your risk committee before launch.

■ Unacceptable

Only through a rule, for prohibited uses such as social scoring. Do not proceed; escalate to your risk committee and legal team.

3. Reading the controls

Controls are grouped by theme: accountability, transparency, data, human oversight, safety and security, and monitoring. Each one says what to do in plain words.

The grey chips are references: the clause, article or principle in each selected framework that asks for the control. ISO/IEC 42001 is cited by clause number only. The numbered links under a control are the official sources used to check it.

Every control carries a verification badge:

✓ **Verified** its references were checked against the cited official sources.

🕒 **Needs verification** the control is sound, but at least one reference has not been checked yet. Hover over the badge to see why, and confirm the reference before you rely on it.

Use the filters to show one framework, or verified controls only.

Under each control, record its status (Met, Partly met, Gap or N/A) and where the evidence is, such as a policy link, a test report or meeting minutes. The tally at the top shows your gaps at a glance. The evidence is saved

with the use case, included in your backup, and printed in the committee report and the XLSX and CSV exports.

For each control, record a status and where the evidence is. It is saved with this use case and included in the committee report and the XLSX and CSV exports.

Evidence so far: 0 Met · 0 Partly met · **0 Gap** · 0 N/A · 53 Not reviewed

Framework ☐ Verified only

Accountability

Approved AI policy ✓ Verified

Adopt an AI policy approved by senior leadership that states the organisation's principles, risk appetite and approval rules for AI, and review it at least once a year.

Dubai: Accountability SDAIA: Accountability & Responsibility ISO 42001: A.2.2, A.2.4 NIST: GOVERN 1.2, GOVERN 2.3

Sources: [\[1\]](#) [\[2\]](#) [\[3\]](#) [\[4\]](#)

Evidence status Evidence

Named accountable owner ⚠ Needs verification

Assign every AI use case a named business owner who is accountable for its outcomes, and record the technical, data and risk contacts alongside them.

UAE: Governance and accountability UAE Gov Code: Agentic Reference p.9 Dubai: Accountability SDAIA: Accountability & Responsibility ISO 42001: A.3.2

NIST: GOVERN 2.1

Sources: [\[1\]](#) [\[2\]](#) [\[3\]](#) [\[4\]](#)

Evidence status Evidence

Controls grouped by theme, with references and verification badges.

4. Recommended autonomy

The three prioritisation questions at the end of the deep-dive (usage, complexity and readiness) recommend how much autonomy the AI should have, from level 1 (full autonomous execution) to level 4 (not suitable for AI yet). The levels are numbered as in the UAE government's AI-assistant priority matrix, so a lower number means more autonomy. They do not change the risk tier.

- **Level 1:** Full autonomous execution
- **Level 2:** Supervised autonomy
- **Level 3:** AI assistance, a person decides
- **Level 4:** Not suitable for AI yet (Not ready for autonomy yet. This is about readiness, not whether AI can help.)

A High tier or a regulated domain never allows level 1: the recommendation stops at level 2, supervised autonomy. If the AI as designed acts more autonomously than recommended, the kit shows a warning: reduce its autonomy or improve readiness first. At level 4 the card lists concrete steps to take first, such as simplifying the process, improving data and naming an owner.

5. Export, share and back up

- **Committee report.** On the Export step, choose *Committee PDF*. Your print dialog opens; choose *Save as PDF*. The one-to-two-page summary has the tier, the reasons, the control checklist and a sign-off block.

- **Control lists.** Download the controls as XLSX or CSV to track them in a spreadsheet. Arabic sheets open right to left.

The screenshot shows the 'GCC AI Governance Kit' interface. At the top, there's a dark green header with the kit's name and navigation links: 'Guide', 'How scoring works', 'Templates', 'About the data', 'GitHub', and a language toggle for 'عربي'. Below the header, the main area has a light beige background. It starts with a 'Register /' link, a 'High risk' status indicator, and the text 'Citizen service triage agent'. A progress bar shows six steps: 'Details', 'Quick check', 'Risk tier', 'Deep-dive', 'Controls', and 'Export' (which is the active step, highlighted in dark green). Below the progress bar, there's a white box titled 'Export' with the instruction 'Take the result into your organisation.' Inside this box are three buttons: 'Committee PDF' (dark green), 'Controls (XLSX)' (light green), and 'Controls (CSV)' (light green). Below these buttons, a note says 'The committee report opens your print dialog. Choose "Save as PDF".' A 'Register' button is located at the bottom right of the white box. At the very bottom of the interface, a footer states 'Your data never leaves this browser. · Code MIT · Content CC BY 4.0'.

The Export step: the committee PDF, and the controls as XLSX or CSV.

- **Back up your register.** On the home page, open *Save or restore* and choose *Save a backup (.json)*. Do this regularly: clearing your browser data deletes the register.
- **Share or move it.** Send the .json file to a colleague. They open the kit, choose *Save or restore*, then *Restore a saved register (.json)*, and can replace or merge their own. For a read-only overview, use *Export register (XLSX)*.
- **Prefer spreadsheets?** The [Templates](#) page has the register, impact assessment and control checklist to fill in offline, plus a vendor AI risk assessment, an AI committee charter and a RACI matrix. You don't need them to use the kit. A filled register template can be brought into the app with *Import a filled register* in the *Save or restore* menu: the tiers are worked out again from the answers.

6. Privacy

The kit has no server. Your use cases are stored only in this browser and are never sent anywhere; you can confirm this in your browser's network tools. Other people cannot see them, and neither can other browsers on your computer. In a private window, or where storage is blocked, the kit warns you that changes will be lost when you close the page, so save a backup.

To remove everything, choose *Clear all data* under the register.

Guidance only, not legal advice. Check obligations with your legal and compliance teams.